WHAT IS INSIDE

An AI agent asks. A person decides. Can a person really stay in control?

I built a small open-source approval gateway for MCP tool calls and tried to break it, three times over. I changed things after approval, hid characters in names, replayed requests, pushed secrets through, and let a real client, Claude Code, walk the whole path.

- 210 scripted checks, 81 documented use cases, every one with what I did and what happened
- Six problems found and fixed, and 22 limits written down instead of hidden
- The permit and the receipt: what a downstream tool really gets
- Government, school, business, IT and plant tools, tried the same way
- A checklist of 13 questions to put to any approval step you buy or build

It is written for the people who have to say yes to agents in their company: security leads, IAM and PAM teams, auditors and the curious.

Everything ran on made-up tools and made-up data. The results show what this gateway did on these cases, on my own machines. They do not prove it secure.



Roshan Trivedi
IAM, PAM and agentic AI security. Fourteen years in identity and privileged access. Bengaluru, India.

FIND ME
roshantrivedi.co.in
linkedin.com/in/roshan-trivedi-ciam-49413a54
roshantrivedi.co.in/mcp-lab-portal/

Who Approved That? ROSHAN TRIVEDI

HUMANGATE LAB

Who approved that?

Putting a human approval step in front of AI agent tool calls, and testing it 210 times.



210 SCRIPTED CHECKS

81 USE CASES

22 LIMITS STATED

WRITTEN BY
Roshan Trivedi
IAM, PAM and agentic AI security